

-

Tutorial

R Tutorial (R-Tutorial.html)

ggplot2

ggplot2 Short Tutorial (ggplot2-Tutorial-With-R.html)

ggplot2 Tutorial 1 - Intro (Complete-Ggplot2-Tutorial-Part1-With-R-Code.html)

ggplot2 Tutorial 2 - Theme (Complete-Ggplot2-Tutorial-Part2-Customizing-Theme-With-R-Code.html)

ggplot2 Tutorial 3 - Masterlist (Top50-Ggplot2-Visualizations-MasterList-R-Code.html)

ggplot2 Quickref (ggplot2-cheatsheet.html)

Foundations

Linear Regression (Linear-Regression.html)

Statistical Tests (Statistical-Tests-in-R.html)

Missing Value Treatment (Missing-Value-Treatment-With-R.html)

Outlier Analysis (Outlier-Treatment-With-R.html)

Feature Selection (Variable-Selection-and-Importance-With-R.html)

Model Selection (Model-Selection-in-R.html)

Logistic Regression (Logistic-Regression-With-R.html)

Advanced Linear Regression (Environments.html)

Advanced Regression Models

Advanced Regression Models (adv-regression-models.html)

Time Series

Time Series Analysis (Time-Series-Analysis-With-R.html)

Time Series Forecasting (Time-Series-Forecasting-With-R.html)

More Time Series Forecasting ([Time-Series-Forecasting-With-R-part2.html](#))

High Performance Computing

Parallel computing ([Parallel-Computing-With-R.html](#))

Strategies to Speedup R code ([Strategies-To-Improve-And-Speedup-R-Code.html](#))

Useful Techniques

Association Mining ([Association-Mining-With-R.html](#))

Multi Dimensional Scaling ([Multi-Dimensional-Scaling-With-R.html](#))

Optimization ([Profiling.html](#))

InformationValue package ([Information-Value-With-R.html](#))

Stay up-to-date. Subscribe!

(https://docs.google.com/forms/d/1xkMYkLNFU9U39Dd8S_2JC0p8B5t6_Yq6zUQjanQQJpY/viewform)

Chat! (<https://docs.google.com/forms/d/13GrkCFcNa-TOIIIQghsz2SIEbc-YqY9eJX02B19I5Ow/viewform>)

Contents

Import Data

1. Random Forest Method
2. Relative Importance
4. MARS
5. Step-wise Regression
6. Boruta
7. Information value and Weight of evidence

Feature Selection Approaches

Finding the most important predictor variables (of features) that explains major part of variance of the response variable is key to identify and build high performing models.

Import Data

For illustrating the various methods, we will use the 'Ozone' data from 'mlbench' package, except for Information value method which is applicable for binary categorical response variables.

```
inputData <- read.csv("http://rstatistics.net/wp-content/uploads/2015/09/ozone1.csv", stringsAsFactors=F)
```

1. Random Forest Method

Random forest can be very effective to find a set of predictors that best explains the variance in the response variable.

```
library(party)
cf1 <- cforest(ozone_reading ~ . , data= inputData, control=cforest_unbiased(mtry=2, ntree=50)) # fit the random forest
varimp(cf1) # get variable importance, based on mean decrease in accuracy
```

#=>	Month	Day_of_month	Day_of_week
#=>	0.689167598	0.115937291	-0.004641633
#=>	pressure_height	Wind_speed	Humidity
#=>	5.519633507	0.125868789	3.474611356
#=>	Temperature_Sandburg	Temperature_ElMonte	Inversion_base_height
#=>	12.878794481	14.175901506	4.276103121
#=>	Pressure_gradient	Inversion_temperature	Visibility
#=>	3.234732558	11.738969777	2.283430842

```
varimp(cf1, conditional=TRUE) # conditional=True, adjusts for correlations between predictors
```

#=>	Month	Day_of_month	Day_of_week
#=>	0.08899435	0.19311805	0.02526252
#=>	pressure_height	Wind_speed	Humidity
#=>	0.35458493	-0.19089686	0.14617239
#=>	Temperature_Sandburg	Temperature_ElMonte	Inversion_base_height
#=>	0.74640367	1.19786882	0.69662788
#=>	Pressure_gradient	Inversion_temperature	Visibility
#=>	0.58295887	0.65507322	0.05380003

```
varimpAUC(cf1) # more robust towards class imbalance.
```

#=>	Month	Day_of_month	Day_of_week
#=>	1.12821259	-0.04079495	0.07800158
#=>	pressure_height	Wind_speed	Humidity
#=>	5.85160593	0.11250973	3.32289714
#=>	Temperature_Sandburg	Temperature_ElMonte	Inversion_base_height
#=>	11.97425093	13.66085973	3.70572939
#=>	Pressure_gradient	Inversion_temperature	Visibility
#=>	3.05169171	11.48762432	2.04145930

2. Relative Importance

Using `calc.relimp {relaimpo}`, the relative importance of variables fed into a `lm` model can be determined as a relative percentage.

```
library(relaimpo)
lmMod <- lm(ozone_reading ~ . , data = inputData) # fit lm() model
relImportance <- calc.relimp(lmMod, type = "lmg", rela = TRUE) # calculate relative importance scaled to 100
sort(relImportance$lmg, decreasing=TRUE) # relative importance
```

#=>	Temperature_ElMonte	Temperature_Sandburg	Inversion_temperature
#=>	0.2297491560	0.2095385438	0.1692950876
#=>	pressure_height	Inversion_base_height	Humidity
#=>	0.1104276154	0.1000912612	0.0833080699
#=>	Visibility	Pressure_gradient	Month
#=>	0.0433277042	0.0320457048	0.0164342902
#=>	Wind_speed	Day_of_month	Day_of_week
#=>	0.0034984964	0.0016927799	0.0005912906


4. MARS

The `earth` package implements variable importance based on Generalized cross validation (GCV), number of subset models the variable occurs (`nsubsets`) and residual sum of squares (RSS).

```
library(earth)
marsModel <- earth(ozone_reading ~ ., data=inputData) # build model
ev <- evimp (marsModel) # estimate variable importance
```

#=>	nsubsets	gcv	rss
#=> Temperature_ElMonte	29	100.0	100.0
#=> Pressure_gradient	28	42.5	48.4
#=> pressure_height	26	30.1	38.1
#=> Month9	25	26.1	34.8
#=> Month5	24	21.9	31.7
#=> Month4	23	19.9	30.0
#=> Month3	22	17.6	28.3
#=> Inversion_base_height	21	14.4	26.1
#=> Month11	19	12.3	24.1
#=> Visibility	18	11.4	23.2
#=> Day_of_month23	14	8.9	19.8
#=> Humidity	13	7.4	18.7
#=> Month6	11	5.1	16.6
#=> Temperature_Sandburg	9	7.0	15.6
#=> Wind_speed	7	5.1	13.4
#=> Month12	6	4.2	12.3
#=> Day_of_month9	3	4.6	9.1
#=> Day_of_week4	2	-3.9	5.9
#=> Day_of_month7-unused	1	-4.7	2.8

```
plot(ev)
```



5. Step-wise Regression

If you have large number of predictors (> 15), split the inputData in chunks of 10 predictors with each chunk holding the responseVar.

```

base.mod <- lm(ozone_reading ~ 1 , data= inputData) # base intercept only model
all.mod <- lm(ozone_reading ~ . , data= inputData) # full model with all predictors
stepMod <- step(base.mod, scope = list(lower = base.mod, upper = all.mod), direction =
"both", trace = 0, steps = 1000) # perform step-wise algorithm
shortlistedVars <- names(unlist(stepMod[[1]])) # get the shortlisted variable.
shortlistedVars <- shortlistedVars[!shortlistedVars %in% "(Intercept)"] # remove interc
ept
print(shortlistedVars)
#=> [1] "Temperature_Sandburg" "Humidity" "Temperature_ElMonte"
#=> [4] "Month" "pressure_height" "Inversion_base_height"

```

The output could includes levels within categorical variables, since 'stepwise' is a linear regression based technique, as seen above.

If you have a large number of predictor variables (100+), the above code may need to be placed in a loop that will run stepwise on sequential chunks of predictors. The shortlisted variables can be accumulated for further analysis towards the end of each iteration. This can be very effective method, if you want to (i) be highly selective about discarding valuable predictor variables. (ii) build multiple models on the response variable.

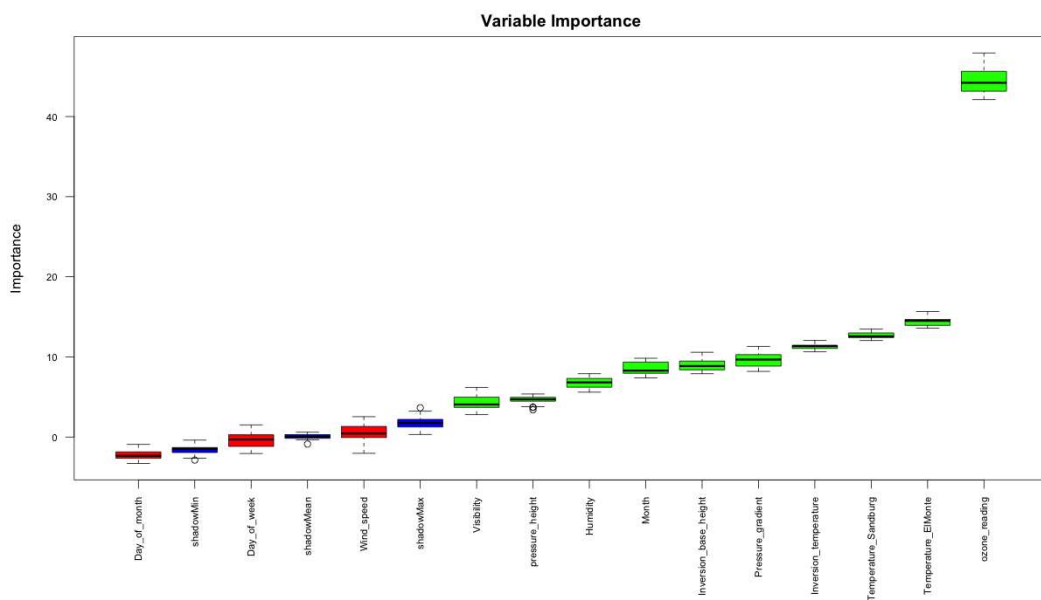
6. Boruta

The 'Boruta' method can be used to decide if a variable is important or not.


```

library(Boruta)
# Decide if a variable is important or not using Boruta
boruta_output <- Boruta(ozone_reading ~ ., data=na.omit(inputData), doTrace=2) # perform Boruta search
# Confirmed 10 attributes: Humidity, Inversion_base_height, Inversion_temperature, Month, Pressure_gradient and 5 more.
# Rejected 3 attributes: Day_of_month, Day_of_week, Wind_speed.
boruta_signif <- names(boruta_output$finalDecision[boruta_output$finalDecision %in% c("Confirmed", "Tentative")]) # collect Confirmed and Tentative variables
print(boruta_signif) # significant variables
#=> [1] "Month" "ozone_reading" "pressure_height"
#=> [4] "Humidity" "Temperature_Sandburg" "Temperature_ElMonte"
#=> [7] "Inversion_base_height" "Pressure_gradient" "Inversion_temperature"
#=> [10] "Visibility"
plot(boruta_output, cex.axis=.7, las=2, xlab="", main="Variable Importance") # plot variable importance

```



7. Information value and Weight of evidence

The InformationValue package (<https://cran.r-project.org/web/packages/InformationValue/vignettes/InformationValue.html>) provides convenient functions to compute *weights of evidence* and *information value* for categorical variables.

Weights of Evidence (WOE) provides a method of recoding a categorical X variable to a continuous variable. For each category of a categorical variable, the **WOE** is calculated as:

$$WOE = \ln \left(\frac{\text{percentage good of all goods}}{\text{percentage bad of all bads}} \right)$$

In above formula, 'goods' is same as 'ones' and 'bads' is same as 'zeros'.

Information Value (IV) is a measure of the predictive capability of a categorical x variable to accurately predict the goods and bads. For each category of x, information value is computed as:

$$InformationValue_{category} = \frac{\text{percentage good of all goods} - \text{percentage bad of all bads}}{WOE}$$

The total IV of a variable is the sum of IV's of its categories.

Example

Let me demonstrate how to create the weights of evidence for categorical variables using the `WOE` function in `InformationValue` pkg. For this we will use the `adult` data as imported below. The response variable in `adult` is the `ABOVE50K` which indicates if the yearly salary of the individual in that row exceeds \$50K. We have a number of predictor variables originally, out of which few of them are categorical variables. On these categorical variables, we will derive the respective `WOEs` using the `InformationValue::WOE` function. Then, let's find out the `InformationValue:IV` of all categorical variables.

Install package from github

```
library(devtools)
install_github("selva86/InformationValue")
```

Import the data

```
library(InformationValue)
inputData <- read.csv("http://rstatistics.net/wp-content/uploads/2015/09/adult.csv")
head(inputData)
```

#=>	AGE	WORKCLASS	FNLWGT	EDUCATION	EDUCATIONNUM	MARITALSTATUS
#=> 1	39	State-gov	77516	Bachelors	13	Never-married
#=> 2	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse
#=> 3	38	Private	215646	HS-grad	9	Divorced
#=> 4	53	Private	234721	11th	7	Married-civ-spouse
#=> 5	28	Private	338409	Bachelors	13	Married-civ-spouse
#=> 6	37	Private	284582	Masters	14	Married-civ-spouse

#	OCCUPATION	RELATIONSHIP	RACE	SEX	CAPITALGAIN	CAPITALLOSS
#=> 1	Adm-clerical	Not-in-family	White	Male	2174	0
#=> 2	Exec-managerial	Husband	White	Male	0	0
#=> 3	Handlers-cleaners	Not-in-family	White	Male	0	0
#=> 4	Handlers-cleaners	Husband	Black	Male	0	0
#=> 5	Prof-specialty	Wife	Black	Female	0	0
#=> 6	Exec-managerial	Wife	White	Female	0	0

#	HOURSPERWEEK	NATIVECOUNTRY	ABOVE50K
#=> 1	40	United-States	0
#=> 2	13	United-States	0
#=> 3	40	United-States	0
#=> 4	40	United-States	0
#=> 5	40	Cuba	0
#=> 6	40	United-States	0

Calculate the Information Values

Below, the information value of each categorical variable is calculated using the `InformationValue::IV` and the strength of each variable is contained within the `howgood` attribute in the returned result. If you are want to dig further into the IV of individual categories within a categorical variable, the `InformationValue::WOETable` (<https://cran.r-project.org/web/packages/InformationValue/vignettes/InformationValue.html#woetable>) will be helpful.

```

factor_vars <- c ("WORKCLASS", "EDUCATION", "MARITALSTATUS", "OCCUPATION", "RELATIONSHIP", "RACE", "SEX", "NATIVECOUNTRY") # get all categorical variables
all_iv <- data.frame(VARS=factor_vars, IV=numeric(length(factor_vars)), STRENGTH=character(length(factor_vars)), stringsAsFactors = F) # init output dataframe
for (factor_var in factor_vars){
  all_iv[all_iv$VARS == factor_var, "IV"] <- InformationValue::IV(X=inputData[, factor_var], Y=inputData$ABOVE50K)
  all_iv[all_iv$VARS == factor_var, "STRENGTH"] <- attr(InformationValue::IV(X=inputData[, factor_var], Y=inputData$ABOVE50K), "howgood")
}

all_iv <- all_iv[order(-all_iv$IV), ] # sort
#>           VARS           IV           STRENGTH
#>  RELATIONSHIP 1.53560810   Highly Predictive
#>  MARITALSTATUS 1.33882907   Highly Predictive
#>    OCCUPATION 0.77622839   Highly Predictive
#>    EDUCATION 0.74105372   Highly Predictive
#>         SEX 0.30328938   Highly Predictive
#>   WORKCLASS 0.16338802   Highly Predictive
#> NATIVECOUNTRY 0.07939344 Somewhat Predictive
#>         RACE 0.06929987 Somewhat Predictive

```

Compute the weights of evidence (optional)

Optionally, we could create the weights of evidence for the factor variables and use it as continuous variables in place of the factors.

```

for(factor_var in factor_vars){
  inputData[[factor_var]] <- WOE(X=inputData[, factor_var], Y=inputData$ABOVE50K)
}
#>  AGE  WORKCLASS FNLWGT  EDUCATION EDUCATIONNUM MARITALSTATUS OCCUPATION
#> 1  39  0.1608547  77516  0.7974104          13    -1.8846680  -0.713645
#> 2  50  0.2254209  83311  0.7974104          13     0.9348331   1.084280
#> 3  38 -0.1278453  215646 -0.5201257           9    -1.0030638  -1.555142
#> 4  53 -0.1278453  234721 -1.7805021           7     0.9348331  -1.555142
#> 5  28 -0.1278453  338409  0.7974104          13     0.9348331   0.943671
#> 6  37 -0.1278453  284582  1.3690863          14     0.9348331   1.084280

#>  RELATIONSHIP      RACE      SEX CAPITALGAIN CAPITALLOSS HOURSPERWEEK
#> 1    -1.015318  0.08064715  0.3281187        2174           0          40
#> 2     0.941801  0.08064715  0.3281187           0           0          13
#> 3    -1.015318  0.08064715  0.3281187           0           0          40
#> 4     0.941801 -0.80794676  0.3281187           0           0          40
#> 5     1.048674 -0.80794676 -0.9480165           0           0          40
#> 6     1.048674  0.08064715 -0.9480165           0           0          40

#>  NATIVECOUNTRY ABOVE50K
#> 1    0.02538318      0
#> 2    0.02538318      0
#> 3    0.02538318      0
#> 4    0.02538318      0
#> 5    0.11671564      0
#> 6    0.02538318      0

```

The newly created woe variables can alternatively be in place of the original factor variables.

© 2016-17 Selva Prabhakaran. Powered by jekyll (<http://jekyllrb.com/>), knitr (<http://yihui.name/knitr/>), and pandoc (<http://johnmacfarlane.net/pandoc/>). This work is licensed under the Creative Commons License. (<http://creativecommons.org/licenses/by-nc/3.0/>)