

Homework_2

Hamed

2/12/2020

1. Load the required packages and the cpus dataset from the MASS package

```
## -- Attaching packages ----- tidyverse
1.3.0 --

## <U+2713> ggplot2 3.2.1      <U+2713> purrr  0.3.3
## <U+2713> tibble  2.1.3      <U+2713> dplyr  0.8.3
## <U+2713> tidyr   1.0.0      <U+2713> stringr 1.4.0
## <U+2713> readr   1.3.1      <U+2713> forcats 0.4.0

## -- Conflicts -----
tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## x dplyr::select() masks MASS::select()

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
##     lift

##           name syct mmin  mmax cach chmin chmax perf estperf
## 1  ADVISOR 32/60  125  256  6000  256   16   128  198    199
## 2  AMDAHL 470V/7   29 8000 32000   32    8    32  269    253
## 3  AMDAHL 470/7A   29 8000 32000   32    8    32  220    253
## 4  AMDAHL 470V/7B  29 8000 32000   32    8    32  172    253
## 5  AMDAHL 470V/7C  29 8000 16000   32    8    16  132    132
## 6  AMDAHL 470V/8   26 8000 32000   64    8    32  318    290
```

2. Use syct, mmin, mmax, cach, chmin, chmax as the predictors (independent variables) to predict performance (perf).

From the model output we use the p-value to check the best predictors in the model, condition: (p-value < 0.05).

```
my_model<-lm(perf~syct + mmin + mmax + cach + chmin + chmax, data = cpus,
nvmax = 6)
summary(my_model)
```

```
##
## Call:
## lm(formula = perf ~ syct + mmin + mmax + cach + chmin + chmax,
##     data = cpus, nvmax = 6)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -195.84  -25.17    5.41   26.53  385.75
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -5.590e+01  8.045e+00  -6.948 4.99e-11 ***
## syct         4.886e-02  1.752e-02   2.789 0.00579 **
## mmin         1.529e-02  1.827e-03   8.371 9.42e-15 ***
## mmax         5.571e-03  6.418e-04   8.680 1.33e-15 ***
## cach         6.412e-01  1.396e-01   4.594 7.64e-06 ***
## chmin        -2.701e-01  8.557e-01  -0.316 0.75263
## chmax         1.483e+00  2.201e-01   6.738 1.64e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 59.99 on 202 degrees of freedom
## Multiple R-squared:  0.8649, Adjusted R-squared:  0.8609
## F-statistic: 215.5 on 6 and 202 DF,  p-value: < 2.2e-16
```

3. Perform best subset selection in order to choose the best predictors from the above predictors. What is the best model obtained according to Cp, BIC, and adjusted R2?

```
models <- regsubsets(perf~syct + mmin + mmax + cach + chmin + chmax,
                     data = cpus, nvmax = 6)
summary(models)

## Subset selection object
## Call: regsubsets.formula(perf ~ syct + mmin + mmax + cach + chmin +
##     chmax, data = cpus, nvmax = 6)
## 6 Variables (and intercept)
##      Forced in Forced out
## syct      FALSE      FALSE
## mmin      FALSE      FALSE
## mmax      FALSE      FALSE
## cach      FALSE      FALSE
## chmin     FALSE      FALSE
## chmax     FALSE      FALSE
## 1 subsets of each size up to 6
## Selection Algorithm: exhaustive
##      syct mmin mmax cach chmin chmax
## 1 ( 1 ) " " " " "*" " " " " " "
## 2 ( 1 ) " " " " "*" "*" " " " "
## 3 ( 1 ) " " "*" "*" " " " " "*"
## 4 ( 1 ) " " "*" "*" "*" " " "*"

```

```
## 5 ( 1 ) "*" "*" "*" "*" " " "*"
## 6 ( 1 ) "*" "*" "*" "*" "*" "*"

res.sum <- summary(models)
```

Displays the CP values at each predictor number from 1-predictor to 6-predictor

```
res.sum$cp
## [1] 176.563616 95.808585 28.225948 10.977588 5.099604 7.000000
```

Displays the BIC values at each predictor number from 1-predictor to 6-predictor

```
res.sum$bic
## [1] -274.7146 -320.4675 -370.5300 -383.5185 -386.1684 -380.9290
```

This shows the number of predictors that BIC, CP, ADJ.R2 support as best subset is the one with lowest BIC, CP, ADJ.R2 values.

```
data.frame(
  Adj.R2 = which.max(res.sum$adjr2),
  CP = which.min(res.sum$cp),
  BIC = which.min(res.sum$bic)
)
##   Adj.R2 CP BIC
## 1      5 5  5
```

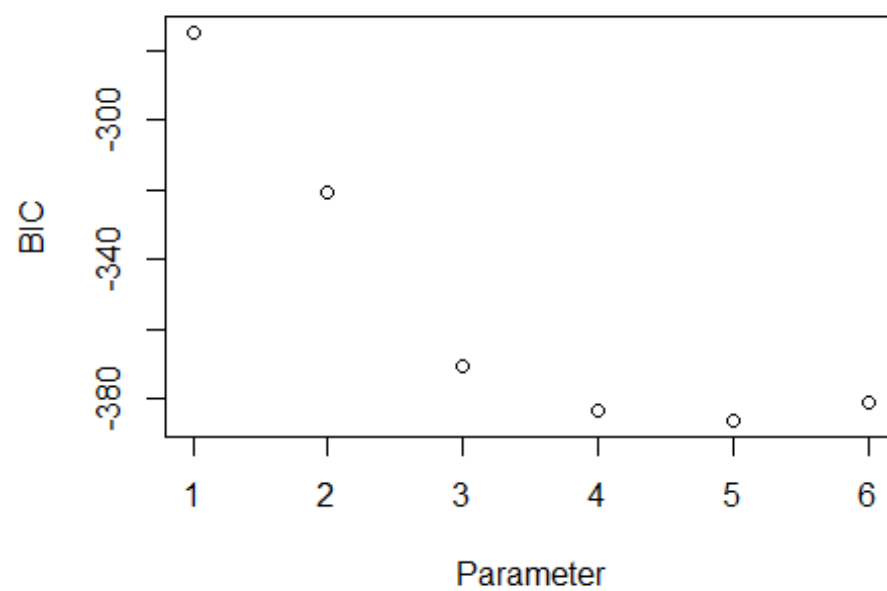
- The function `summary()` reports the best set of variables for each model size. From the output, an asterisk specifies that a given variable is included in the corresponding model.
- For example, it can be seen that the best 2-variables model contains only `mmax` and `cach` variables (`perf ~ mmax + cach`). The best three-variable model is (`perf ~ mmax + cach + mmin`), and so forth.
- As shown above, adjusted R2, BIC and Cp criteria, tells us that the best model is the one with 5 predictor variables.
- A natural question is: which of these best models should we finally choose for our predictive analytics?

4. Show some plots to provide evidence for your answer, and report the coefficients of the best model obtained for each criteria.

- The plots show that the BIC reduces as the number of parameters increase upto 5 then it becomes constant. The same goes for CP.

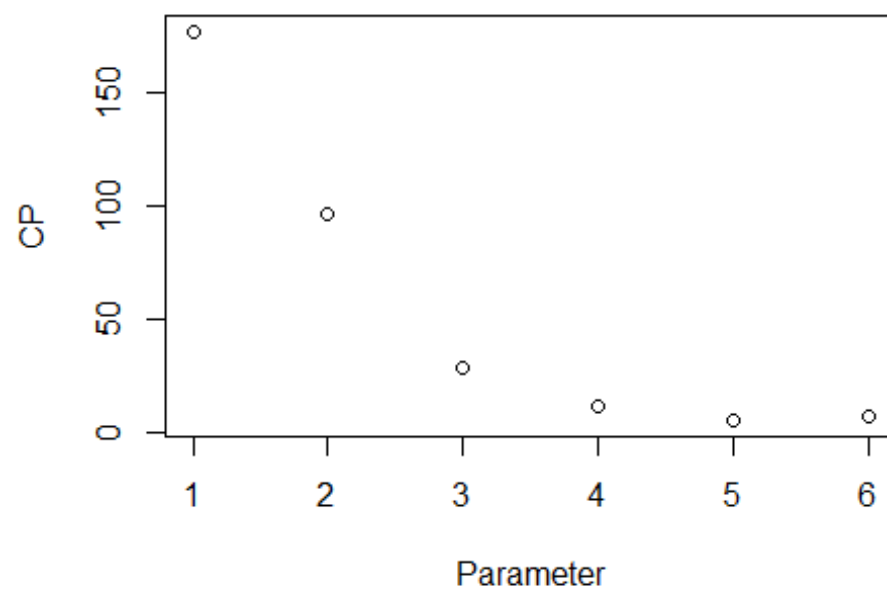
```
plot(res.sum$bic, xlab="Parameter", ylab="BIC", main="BIC plot")
```

BIC plot



```
plot(res.sum$cp, xlab="Parameter", ylab="CP",main = "CP plot")
```

CP plot



5. Repeat using forward stepwise selection and also using backwards stepwise selection. How does your answer compare to the best subset results?

- nvmax: the number of variable in the model. For example nvmax = 2, specify the best 2-variables model
- RMSE and MAE are two different metrics measuring the prediction error of each model. The lower the RMSE and MAE, the better the model.
- Rsquared indicates the correlation between the observed outcome values and the values predicted by the model. The higher the R squared, the better the model.

Fit the full model to show the performance before subset selection

```
library(MASS)
full.model <- lm(perf ~ syct + mmin + mmax + cach + chmin + chmax, data =
cpus)
full.model

##
## Call:
## lm(formula = perf ~ syct + mmin + mmax + cach + chmin + chmax,
##     data = cpus)
##
## Coefficients:
## (Intercept)          syct          mmin          mmax          cach
chmin
## -55.900116      0.048863      0.015294      0.005571      0.641207      -
0.270065
##          chmax
##          1.482694
```

Model information

- Specify the tuning parameter nvmax, which corresponds to the maximum number of predictors to be incorporated in the model.
- For example, you can vary nvmax from 1 to 5. In this case, the function starts by searching different best models of different size, up to the best 5-variables model.
- That is, it searches the best 1-variable model, the best 2-variables model, ..., the best 5-variables models.
- As the data set contains only 6 predictors, we'll vary nvmax from 1 to 6 resulting to the identification of the 6 best models with different sizes: the best 1-variable model, the best 2-variables model, ..., the best 6-variables model.
- We'll use 10-fold cross-validation to estimate the average prediction error (RMSE) of each of the 6 models
- The output of the final model has predictors selected based on the number of asterics , the more the better.

Forward selection

```
# Set seed for reproducibility
set.seed(123)
```

```
# Set up repeated k-fold cross-validation
train.control <- trainControl(method = "cv", number = 10)
# Train the model

step.model <- train(perf ~syst + mmin + mmax + cach + chmin + chmax, data =
cpus,
                    method = "leapForward",
                    tuneGrid = data.frame(nvmax = 1:6),
                    trControl = train.control
)
step.model$results
```

##	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	93.46400	0.6488552	56.59877	60.38499	0.1779760	22.54073
## 2	2	89.08646	0.7168064	53.65864	59.70345	0.1501266	21.52435
## 3	3	83.01186	0.7670026	48.96523	59.63821	0.1347547	21.51615
## 4	4	74.41985	0.8004094	45.31268	49.53468	0.1247854	17.23585
## 5	5	72.35948	0.8086084	44.53950	48.65457	0.1051621	17.19711
## 6	6	72.78049	0.8063649	44.82444	47.12075	0.1081818	16.66283

Here we show the number of predictors the best subset will have after forward selection

```
step.model$bestTune
```

```
##   nvmax
## 5     5
```

The final model after forward selection showing the importance of each predictor.

```
summary(step.model$finalModel)
```

```
## Subset selection object
## 6 Variables (and intercept)
##      Forced in Forced out
## syst      FALSE      FALSE
## mmin      FALSE      FALSE
## mmax      FALSE      FALSE
## cach      FALSE      FALSE
## chmin     FALSE      FALSE
## chmax     FALSE      FALSE
## 1 subsets of each size up to 5
## Selection Algorithm: forward
##      syst mmin mmax cach chmin chmax
## 1 ( 1 ) " " " " "*" " " " " " "
## 2 ( 1 ) " " " " "*" "*" " " " "
## 3 ( 1 ) " " "*" "*" "*" " " " "
## 4 ( 1 ) " " "*" "*" "*" " " "*"
## 5 ( 1 ) "*" "*" "*" "*" " " "*"
```

```
set.seed(123)
```

```
# Set up repeated k-fold cross-validation
train.control <- trainControl(method = "cv", number = 10)
# Train the model

# Train the model
step.model2 <- train(perf ~syst + mmin + mmax + cach + chmin + chmax, data =
cpus,
                    method = "leapBackward",
                    tuneGrid = data.frame(nvmax = 1:6),
                    trControl = train.control)
step.model2$results
```

##	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	90.67875	0.6109622	53.41392	62.55296	0.2355191	24.00383
## 2	2	88.89009	0.7007681	52.00267	62.55203	0.1235651	23.38322
## 3	3	76.09071	0.7765757	46.44309	60.69159	0.1381791	23.17784
## 4	4	74.41985	0.8004094	45.31268	49.53468	0.1247854	17.23585
## 5	5	72.35948	0.8086084	44.53950	48.65457	0.1051621	17.19711
## 6	6	72.78049	0.8063649	44.82444	47.12075	0.1081818	16.66283

Here we show the number of predictors the best subset will have after Backward selection.

```
step.model2$bestTune
```

```
##      nvmax
## 5      5
```

The final model after after Backward selection showing the importance of each predictor.

```
summary(step.model2$finalModel)

## Subset selection object
## 6 Variables (and intercept)
##           Forced in Forced out
## syct      FALSE      FALSE
## mmin      FALSE      FALSE
## mmax      FALSE      FALSE
## cach      FALSE      FALSE
## chmin     FALSE      FALSE
## chmax     FALSE      FALSE
## 1 subsets of each size up to 5
## Selection Algorithm: backward
##           syct mmin mmax cach chmin chmax
## 1  ( 1 ) " "  "*"  " "  " "  " "  " "
## 2  ( 1 ) " "  "*"  " "  " "  " "  "*"
## 3  ( 1 ) " "  "*"  "*"  " "  " "  "*"
## 4  ( 1 ) " "  "*"  "*"  "*"  " "  "*"
## 5  ( 1 ) "*"  "*"  "*"  "*"  " "  "*"

```

Stepwise selection

```
set.seed(123)
# Set up repeated k-fold cross-validation
train.control <- trainControl(method = "cv", number = 10)
# Train the model

step.model3 <- train(perf ~sycr + mmin + mmax + cach + chmin + chmax, data =
cpus,
                      method = "leapSeq",
                      tuneGrid = data.frame(nvmax = 1:6),
                      trControl = train.control)
step.model3$results
```

##	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	93.46400	0.6488552	56.59877	60.38499	0.17797603	22.54073
## 2	2	88.10355	0.7235298	52.90480	60.02006	0.14462287	21.58357
## 3	3	76.85737	0.7660263	46.67915	60.69824	0.15230524	23.38628
## 4	4	73.30065	0.8123798	45.88578	48.41967	0.08509873	16.35415
## 5	5	72.00203	0.8281705	44.90645	47.86131	0.08199148	16.33195
## 6	6	72.78049	0.8063649	44.82444	47.12075	0.10818179	16.66283

Here we show the number of predictors the best subset will have

```
step.model3$bestTune
```

```
## nvmax
## 5      5
```

The final model after stepwise selection showing the importance of each predictor.

```
summary(step.model3$finalModel)
```

```
## Subset selection object
## 6 Variables (and intercept)
##      Forced in Forced out
## sycr      FALSE      FALSE
## mmin      FALSE      FALSE
## mmax      FALSE      FALSE
## cach      FALSE      FALSE
## chmin     FALSE      FALSE
## chmax     FALSE      FALSE
## 1 subsets of each size up to 5
## Selection Algorithm: 'sequential replacement'
##      sycr mmin mmax cach chmin chmax
## 1 ( 1 ) " " " " "*" " " " " " "
## 2 ( 1 ) " " " " "*" "*" " " " "
## 3 ( 1 ) " " "*" "*" " " " " "*"
## 4 ( 1 ) "*" "*" "*" "*" " " " "
## 5 ( 1 ) "*" "*" "*" "*" "*" " "
```


REGRESSION METHODS: OLS, RIDGE, LASSO, PCR, & PLS USING “College” dataset

Predict the number of applications received using the other variables in the College data set in library ISLR

(a) Split the data set into a training set and a test set using caret library and fit each of the following models using caret and ten fold cross validation.

```
library(ISLR)
library(glmnet)
attach(College)
head(College)
```

```
##                               Private Apps Accept Enroll Top10perc
Top25perc
## Abilene Christian University   Yes 1660   1232    721         23
52
## Adelphi University            Yes 2186   1924    512         16
29
## Adrian College                Yes 1428   1097    336         22
50
## Agnes Scott College           Yes  417    349     137         60
89
## Alaska Pacific University     Yes  193    146      55         16
44
## Albertson College             Yes  587    479     158         38
62
##                               F.Undergrad P.Undergrad Outstate Room.Board
Books
## Abilene Christian University    2885         537    7440        3300
450
## Adelphi University             2683         1227   12280        6450
750
## Adrian College                 1036          99   11250        3750
400
## Agnes Scott College            510          63   12960        5450
450
## Alaska Pacific University       249          869    7560        4120
800
## Albertson College              678          41   13500        3335
500
##                               Personal PhD Terminal S.F.Ratio perc.alumni
Expend
## Abilene Christian University   2200   70      78     18.1         12
7041
## Adelphi University            1500   29      30     12.2         16
10527
## Adrian College               1165   53      66     12.9         30
8735
## Agnes Scott College           875   92      97      7.7         37
19016
```

```
## Alaska Pacific University      1500  76      72      11.9      2
10922
## Albertson College              675  67      73      9.4      11
9727
##                               Grad.Rate
## Abilene Christian University    60
## Adelphi University             56
## Adrian College                 54
## Agnes Scott College            59
## Alaska Pacific University       15
## Albertson College              55

x <- model.matrix(Apps~., College)[-1]
y <- College$Apps
lambda <- 10^seq(10, -2, length = 100)

# Train test split
set.seed(489)
train = sample(1:nrow(x), nrow(x)/2)
test = (-train)
ytest = y[test]
```

(b) Fit a linear model using ordinary least squares on the training set, and report the test mean squared error obtained.

```
OLS_lm <- lm(Apps~., data = College, subset = train)
OLS_lm

##
## Call:
## lm(formula = Apps ~ ., data = College, subset = train)
##
## Coefficients:
## (Intercept)  PrivateYes      Accept      Enroll  Top10perc
Top25perc
## -544.41744   -170.52279     1.74160    -1.41087    38.28257    -
6.06587
## F.Undergrad  P.Undergrad  Outstate  Room.Board      Books
Personal
## 0.07306      0.08748    -0.08632     0.16650     0.06319
0.09351
##          PhD      Terminal  S.F.Ratio  perc.alumni      Expend
Grad.Rate
## -11.10782    2.19668     4.12585     3.56206     0.05095
1.92934

#Find the best lambda from our list via cross-validation
cv.out <- cv.glmnet(x[train,], y[train], alpha = 0)
cv.out
```

```
##
## Call:  cv.glmnet(x = x[train, ], y = y[train], alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Measure      SE Nonzero
## min   397.4 2103455 1270039      17
## 1se 2554.6 3360297 2169940      17

#Best Lambda
bestlam <- cv.out$lambda.min
bestlam

## [1] 397.4201

#Make predictions
OLS.pred <- predict(OLS_lm, newdata = College[test,])
head(OLS.pred)

##      Adelphi University      Adrian College      Albertson
College
##      3350.61158      1397.93516
608.67123
##      Albertus Magnus College Alderson-Broadus College      Allegheny
College
##      54.98646      686.22811
2922.74735
```

- MEAN SQUARED ERROR FOR OLS

```
#check Mean Squared Error
mean((OLS.pred-ytest)^2)

## [1] 1403054
```

(c) Fit a ridge regression model on the training set, with λ chosen by cross-validation. Report the test mean squared error obtained. Report the value of λ used in the model

```
ridge.mod <- glmnet(x[train,], y[train], alpha = 0, lambda = lambda)
summary(ridge.mod)

##      Length Class      Mode
## a0      100   -none-   numeric
## beta    1700  dgMatrix S4
## df      100   -none-   numeric
## dim      2    -none-   numeric
## lambda   100   -none-   numeric
## dev.ratio 100   -none-   numeric
## nulldev    1   -none-   numeric
## npasses    1   -none-   numeric
## jerr       1   -none-   numeric
## offset     1   -none-   logical
```

```
## call          5  -none-    call
## nobs          1  -none-    numeric

#Find the best lambda from our list via cross-validation
cv.out <- cv.glmnet(x[train,], y[train], alpha = 0)
cv.out

##
## Call:  cv.glmnet(x = x[train, ], y = y[train], alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Measure      SE Nonzero
## min  397.4 2352967 1646036      17
## 1se 3077.1 3903384 2937840      17

#Best Lambda
bestlam <- cv.out$lambda.min
bestlam

## [1] 397.4201

#make predictions
ridge.pred <- predict(ridge.mod, s = bestlam, newx = x[test,])
head(ridge.pred)

##                                1
## Adelphi University      3000.9738
## Adrian College          1164.0138
## Albertson College        595.0114
## Albertus Magnus College  317.8752
## Alderson-Broadbudd College 549.4096
## Allegheny College       2677.7668
```

- MEAN SQUARED ERROR FOR RIDGE REGRESSION

```
#Mean squared error
mean((ridge.pred-ytest)^2)

## [1] 1298095
```

(d) Fit a lasso model on the training set, with fraction chosen by cross validation. Report the test mean squared error obtained, along with the number of non-zero coefficient estimates and the fraction.

```
lasso.mod <- glmnet(x[train,], y[train], alpha = 1, lambda = lambda)
summary(lasso.mod)

##      Length Class      Mode
## a0      100  -none-    numeric
## beta    1700 dgCMatrix S4
## df       100  -none-    numeric
## dim        2  -none-    numeric
```

```
## lambda      100    -none-    numeric
## dev.ratio   100    -none-    numeric
## nulldev     1     -none-    numeric
## npasses     1     -none-    numeric
## jerr        1     -none-    numeric
## offset      1     -none-    logical
## call        5     -none-    call
## nobs        1     -none-    numeric
```

```
lasso.pred <- predict(lasso.mod, s = bestlam, newx = x[test,])
head(lasso.pred)
```

```
##                                1
## Adelphi University           2741.3266
## Adrian College              1686.0656
## Albertson College            998.9299
## Albertus Magnus College      629.1303
## Alderson-Broadbudd College   875.6115
## Allegheny College            2954.1927
```

- MEAN SQUARED ERROR FOR LASSO REGRESSION

```
mean((lasso.pred-ytest)^2)
```

```
## [1] 1798354
```

(e) Fit a PCR model on the training set, with no. of principal components M chosen by cross-validation. Report the test mean squared error obtained, along with the value of M selected by cross-validation.

```
set.seed(123)
smp_size <- floor(0.75 * nrow(mtcars))
train_ind <- sample(seq_len(nrow(College)), size = smp_size)
train_p <- College[train_ind, ]
test_p <- College[-train_ind,c(1,3:18) ]
y_test=College[-train_ind,2]

require(pls)
pcr_model <- pcr(Apps~., data = train_p,scale =TRUE, validation = "CV")
summary(pcr_model)
```

```
## Data:      X dimension: 24 17
## Y dimension: 24 1
## Fit method: svdpc
## Number of components considered: 17
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              2426    2779    1375    1389    1371    1509    1612
## adjCV           2426    2749    1351    1365    1342    1477    1568
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
```

```
## CV      1605      1625      1842      1786      1664      1373      1346
## adjCV    1559      1574      1779      1714      1606      1305      1282
##      14 comps  15 comps  16 comps  17 comps
## CV      1180      936.9      1293      2503
## adjCV    1126      890.5      1224      2386
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8
comps
## X      37.40    62.89    73.63    81.16    87.87    91.82    93.96
95.65
## Apps   23.12    81.60    82.53    84.39    86.72    89.36    90.93
91.92
##      9 comps  10 comps  11 comps  12 comps  13 comps  14 comps  15 comps
## X      97.14    97.93    98.65    99.09    99.50    99.79    99.96
## Apps   92.87    95.00    95.49    98.20    98.27    98.44    98.99
##      16 comps  17 comps
## X      99.98    100.00
## Apps   98.99    99.03

pcr_pred <- predict(pcr_model, test_p, ncomp = 3)
head(pcr_pred)

## [1] 1930.6961 1451.1950 704.8014 2322.1893 815.2842 1231.9749
```

- MEAN SQUARED ERROR FOR Principal component regression (PCR)

```
mean((pcr_pred - y_test)^2)
```

```
## [1] 3664827
```

(f) Fit a PLS model on the training set, with M chosen by cross validation. Report the test error obtained, along with the value of M selected by cross-validation.

```
library(caret)
# Compile cross-validation settings
set.seed(100)
myfolds <- createMultiFolds(train_p$Apps, k = 5, times = 10)
control <- trainControl("repeatedcv", index = myfolds, selectionFunction =
"oneSE")

# Train PLS model
mod1 <- train(Apps ~ ., data = train_p,
              method = "pls",
              metric = "RMSE",
              tuneLength = 20,
              trControl = control,
              preProc = c("zv", "center", "scale"))

summary(mod1)
```

```
## Data:      X dimension: 24 17
## Y dimension: 24 1
## Fit method: oscorespls
## Number of components considered: 8
## TRAINING: % variance explained
##           1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X           30.84   61.11   69.40   75.50   81.94   85.84   89.89
## .outcome    82.50   87.84   92.86   95.44   96.89   97.95   98.46
##           8 comps
## X           94.22
## .outcome    98.71
```

The model results display the metrics in the model including: ncom (number of predictors in a subset, also is the value of M), root mean squared error (RMSE), R-squared, mean absolute error (MAE) etc. The lowest RMSE indicates the best subset size and return the best model.

```
mod1$results
```

##	ncomp	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	1084.9854	0.7133798	861.7848	521.4875	0.3403221	332.2876
## 2	2	1180.7783	0.7798500	875.6709	564.5886	0.3116584	354.9473
## 3	3	1230.5371	0.7154510	897.1544	561.0249	0.3305893	389.9395
## 4	4	1191.8641	0.7283764	905.3304	475.5401	0.2813624	375.5198
## 5	5	1132.9057	0.7626434	881.3255	443.3958	0.2641401	354.1532
## 6	6	1074.9639	0.7841007	846.2498	384.8527	0.2348362	305.2852
## 7	7	1030.9962	0.8037494	823.7655	344.5840	0.2158111	280.5208
## 8	8	977.8485	0.8270842	799.8415	320.7414	0.1946075	256.7419
## 9	9	961.4021	0.8474581	796.9018	369.8159	0.1860991	279.0550
## 10	10	1003.3505	0.8466150	833.7763	403.9196	0.1928643	310.0666
## 11	11	1066.8502	0.8373313	887.4662	421.3553	0.1989530	320.6783
## 12	12	1137.5497	0.8239245	944.5730	463.5097	0.2021272	351.6782
## 13	13	1229.7565	0.7731685	1007.7498	525.0526	0.2575168	399.2797
## 14	14	1472.4609	0.7342725	1172.4776	702.0939	0.2797465	505.0709
## 15	15	1936.8054	0.6771353	1474.1501	1218.4259	0.3113862	797.1762
## 16	16	2309.6684	0.6659932	1770.6816	1720.6389	0.3131435	1165.2617

(g) Comment on the results obtained. Is there much difference among the test errors resulting from these five approaches?

- There is a noticeable difference between OLS, Ridge, PCR and PLS regression in terms of mean squared error whereby Ridge regression had the lowest mean squared error followed by PLS, OLS, Lasso and then Principal Component Regression had the highest mean squared error.