



R news and tutorials contributed by hundreds of R bloggers

- [Home](#)
- [About](#)
- [RSS](#)
- [add your blog!](#)
- [Learn R](#)
- [R jobs](#)
- [Contact us](#)

Welcome!

Follow @rbloggers

78.9k

Here you will find daily news and tutorials about R, contributed by hundreds of bloggers.

There are many ways to follow us -

[By e-mail:](#)

51572 readers

BY FEEDBURNER

[On Facebook:](#)

R blog...

77K likes

Like Page

3 friends like this

If you are an R blogger yourself you are invited to [add your own R content feed to this site](#) (Non-English R bloggers should add themselves- [here](#))

[Jobs for R-users](#)

- [\(Junior\) Data Analyst – Berlin](#)
- [Associate Researcher \(Data Analytics\)](#)
- [Lead Data Scientist @ Stabio, Ticino, Switzerland](#)
- [Senior Scientist, Translational Informatics](#)
- [Senior Scientist, Translational Informatics @](#)

Recent Posts

- [dang 0.0.11: Small improvements](#)
- [The politics of New Mexico: a brief historical-visual account](#)
- [RStudio Professional Drivers 1.6.0](#)
- [Gold-Mining Week 8 \(2019\)](#)
- [Decision Making Support Systems #3: Differences between IA and AI](#)
- [Horizontal scaling of data science applications in the cloud](#)
- [lml 0.0.4: Now with footer](#)
- [pkgKitten 0.1.5: Creating R Packages that purr](#)
- [Council spending – open data](#)
- [Super Solutions for Shiny Architecture #5 of 5: Automated Tests](#)
- [Understanding Blockchain Technology by building one in R](#)
- [digest 0.6.22: More goodies!](#)
- [Ordering Sentinel-2 products from Long Term Archive with sen2r](#)
- [Access the free economic database DBnomics with R](#)
- [Trends in U.S. Border Crossing Entry since 1996](#)

Other sites

- [Jobs for R-users](#)
- [SAS blogs](#)

Fitting a Model by Maximum Likelihood

August 18, 2013

By [andrew](#)

Like 9 Share Tweet Share

[This article was first published on [Exegetic Analytics » R](#), and kindly contributed to [R-bloggers](#). (You can report issue about the content on this page [here](#))

I Want to share your content on R-bloggers? [click here](#) if you have a blog or [here](#) if you don't

[Share](#)[Tweet](#)

Maximum-Likelihood Estimation (MLE) is a statistical technique for estimating model parameters. It basically sets out to answer the question: what model parameters are most likely to characterise a given set of data? First you need to select a model for the data. And the model must have one or more (unknown) parameters. As the name implies, MLE proceeds to maximise a likelihood function, which in turn maximises the agreement between the model and the data.

Most illustrative examples of MLE aim to derive the parameters for a probability density function (PDF) of a particular distribution. In this case the likelihood function is obtained by considering the PDF not as a function of the sample variable, but as a function of distribution's parameters. For each data point one then has a function of the distribution's parameters. The joint likelihood of the full data set is the product of these functions. This product is generally very small indeed, so the likelihood function is normally replaced by a log-likelihood function. Maximising either the likelihood or log-likelihood function yields the same results, but the latter is just a little more tractable!

Fitting a Normal Distribution

Let's illustrate with a simple example: fitting a normal distribution. First we generate some data.

```
> set.seed(1001)
>
> N <- 100
>
> x <- rnorm(N, mean = 3, sd = 2)
>
> mean(x)
[1] 2.998305
> sd(x)
[1] 2.288979
```

Then we formulate the log-likelihood function.

```
> LL <- function(mu, sigma) {
+   R = dnorm(x, mu, sigma)
+   #
+   -sum(log(R))
+ }
```

And apply MLE to estimate the two parameters (mean and standard deviation) for which the normal distribution best describes the data.

```
> library(stats4)
>
> mle(LL, start = list(mu = 1, sigma=1))

Call:
mle(minuslogl = LL, start = list(mu = 1, sigma = 1))

Coefficients:
      mu      sigma 
2.998305 2.277506 

Warning messages:
1: In dnorm(x, mu, sigma) : NaNs produced
2: In dnorm(x, mu, sigma) : NaNs produced
3: In dnorm(x, mu, sigma) : NaNs produced
```

Those warnings are a little disconcerting! They are produced when negative values are attempted for the standard deviation.

```
> dnorm(x, 1, -1)
[1] NaN NaN
[30] NaN NaN
[59] NaN NaN
[88] NaN NaN
```

There are two ways to sort this out. The first is to apply constraints on the parameters. The mean does not require a constraint but we insist that the standard deviation is positive.

```
> mle(LL, start = list(mu = 1, sigma=1), method = "L
  upper = c(Inf, Inf))

Call:
mle(minuslogl = LL, start = list(mu = 1, sigma = 1),
lower = c(-Inf, 0), upper = c(Inf, Inf))

Coefficients:
      mu     sigma
2.998304 2.277506
```

This works because `mle()` calls `optim()`, which has a number of optimisation methods. The default method is BFGS. An alternative, the L-BFGS-B method, allows box constraints.

The other solution is to simply ignore the warnings. It's neater and produces the same results.

```
> LL <- function(mu, sigma) {
+   R = suppressWarnings(dnorm(x, mu, sigma))
+   #
+   -sum(log(R))
+ }
>
> mle(LL, start = list(mu = 1, sigma=1))

Call:
mle(minuslogl = LL, start = list(mu = 1, sigma = 1))

Coefficients:
      mu     sigma
2.998305 2.277506
```

The maximum-likelihood values for the mean and standard deviation are damn close to the corresponding sample statistics for the data. Of course, they do not agree perfectly with the values used when we generated the data: the results can only be as good as the data. If there were more samples then the results would be closer to these ideal values.

A note of caution: if your initial guess for the parameters is too far off then things can go seriously wrong!

```
> mle(LL, start = list(mu = 0, sigma=1))

Call:
mle(minuslogl = LL, start = list(mu = 0, sigma = 1))

Coefficients:
      mu     sigma
51.4840 226.8299
```

Fitting a Linear Model

Now let's try something a little more sophisticated: fitting a linear model. As before, we generate some data.

```
> x <- runif(N)
> y <- 5 * x + 3 + rnorm(N)
```

We can immediately fit this model using least squares regression.

```
> fit <- lm(y ~ x)
>
> summary(fit)

Call:
```

```
Residuals:
    Min      1Q  Median      3Q     Max 
-1.96206 -0.59016 -0.00166  0.51813  2.43778 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 3.1080    0.1695   18.34 <2e-16 ***
x           4.9516    0.2962   16.72 <2e-16 ***  
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 

Residual standard error: 0.8871 on 98 degrees of freedom
Multiple R-squared:  0.7404, Adjusted R-squared:
F-statistic: 279.5 on 1 and 98 DF, p-value: < 2.2e-
```

The values for the slope and intercept are very satisfactory. No arguments there. We can superimpose the fitted line onto a scatter plot.

```
> plot(x, y)
> abline(fit, col = "red")
```



Pushing on to the MLE for the linear model parameters. First we need a likelihood function. The model is not a PDF, so we can't proceed in precisely the same way that we did with the normal distribution. However, if you fit a linear model then you want the residuals to be normally distributed. So the likelihood function fits a normal distribution to the residuals.

```
LL <- function(beta0, beta1, mu, sigma) {
  # Find residuals
  #
  R = y - x * beta1 - beta0
  #
  # Calculate the likelihood for the residuals (wi
  #
  R = suppressWarnings(dnorm(R, mu, sigma))
  #
  # Sum the log likelihoods for all of the data po
  #
  -sum(log(R))
}
```

One small refinement that one might make is to move the logarithm into the call to dnorm().

```
LL <- function(beta0, beta1, mu, sigma) {
  R = y - x * beta1 - beta0
  #
  R = suppressWarnings(dnorm(R, mu, sigma, log = T
  #
  -sum(R)
}
```

It turns out that the initial guess is again rather important and a poor choice can result in errors. We will return to this issue a little later.

```
> fit <- mle(LL, start = list(beta0 = 3, beta1 = 1,
Error in solve.default(oout$hessian) :
  system is computationally singular: reciprocal con
> fit <- mle(LL, start = list(beta0 = 5, beta1 = 3,
Error in solve.default(oout$hessian) :
  Lapack routine dgesv: system is exactly singular:
```

But if we choose values that are reasonably close then we get a decent outcome.

```
> fit <- mle(LL, start = list(beta0 = 4, beta1 = 2,
> fit
>
> call:
>
```

```
Coefficients:
      beta0      beta1        mu      sigma
 3.5540217  4.9516133 -0.4459783  0.8782272
```

The maximum-likelihood estimates for the slope (beta1) and intercept (beta0) are not too bad. But there is a troubling warning about NaNs being produced in the summary output below.

```
> summary(fit)
Maximum likelihood estimation

Call:
mle(minuslogl = LL, start = list(beta0 = 4, beta1 =
sigma = 1))

Coefficients:
            Estimate Std. Error
beta0    3.5540217   NaN
beta1    4.9516133  0.2931924
mu       -0.4459783   NaN
sigma    0.8782272  0.0620997

-2 log L: 257.8177
Warning message:
In sqrt(diag(object@vcov)) : NaNs produced
```

It stands to reason that we actually want to have the zero mean for the residuals. We can apply this constraint by specifying mu as a fixed parameter. Another option would be to simply replace mu with 0 in the call to dnorm(), but the alternative is just a little more flexible.

```
> fit <- mle(LL, start = list(beta0 = 2, beta1 = 1.5
  nobs = length(y)))
> summary(fit)
Maximum likelihood estimation

Call:
mle(minuslogl = LL, start = list(beta0 = 2, beta1 =
  fixed = list(mu = 0), nobs = length(y)))

Coefficients:
            Estimate Std. Error
beta0    3.1080423 0.16779428
beta1    4.9516164 0.29319233
sigma    0.8782272 0.06209969

-2 log L: 257.8177
```

The resulting estimates for the slope and intercept are rather good. And we have standard errors for these parameters as well.

How about assessing the overall quality of the model? We can look at the [Akaike Information Criterion](#) (AIC) and [Bayesian Information Criterion](#) (BIC). These can be used to compare the performance of different models for a given set of data.

```
> AIC(fit)
[1] 263.8177
> BIC(fit)
[1] 271.6332
> logLik(fit)
'log Lik.' -128.9088 (df=3)
```

Returning now to the errors mentioned above. Both of the cases where the call to mle() failed resulted from problems with inverting the [Hessian Matrix](#). With the implementation of mle() in the stats4 package there is really no way to get around this problem apart from having a good initial guess. In some situations though, this is just not feasible. There are, however, alternative implementations of MLE which circumvent this problem. The bbmle package has mle2() which offers essentially the same functionality but includes the option of not inverting the Hessian Matrix.

```

>
> fit <- mle2(LL, start = list(beta0 = 3, beta1 = 1,
>
> summary(fit)
Maximum likelihood estimation

Call:
mle2(minuslogl = LL, start = list(beta0 = 3, beta1 =
  sigma = 1))

Coefficients:
            Estimate Std. Error z value Pr(z)
beta0    3.054021   0.083897 36.4019 <2e-16 ***
beta1    4.951617   0.293193 16.8886 <2e-16 ***
mu       0.054021   0.083897  0.6439 0.5196
sigma   0.878228   0.062100 14.1421 <2e-16 ***
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’

-2 log L: 257.8177

```

Here mle2() is called with the same initial guess that broke mle(), but it works fine. The summary information for the optimal set of parameters is also more extensive.

Fitting a linear model is just a toy example. However, Maximum Likelihood Estimation can be applied to models of arbitrary complexity. If the model residuals are expected to be normally distributed then a log-likelihood function based on the one above can be used. If the residuals conform to a different distribution then the appropriate density function should be used instead of dnorm().

 Share

 Tweet

To leave a comment for the author, please follow the link and comment on their blog:
[Exegetic Analytics » R](#)

R-bloggers.com offers [daily e-mail updates](#) about R news and tutorials about [learning R](#) and many other topics. [Click here if you're looking to post or find an R/data-science job](#).

Want to share your content on R-bloggers? [click here](#) if you have a blog, or [here](#) if you don't.

If you got this far, why not [subscribe for updates](#) from the site?
 Choose your flavor: [e-mail](#), [twitter](#), [RSS](#), or [facebook](#)...

Like 9 Share Tweet Share

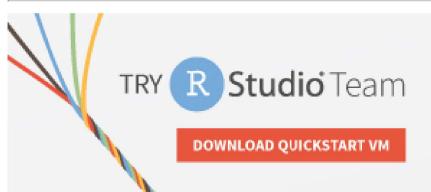
Comments are closed.

Search R-bloggers

Most visited articles of the week

1. [How to write the first for loop in R](#)
2. [5 Ways to Subset a Data Frame in R](#)
3. [R – Sorting a data frame by the contents of a column](#)
4. [Installing R packages](#)
5. [Using apply, sapply, lapply in R](#)
6. [Date Formats in R](#)
7. [How to perform a Logistic Regression in R](#)
8. [RStudio Professional Drivers 1.6.0](#)
9. [R Tutorial Series: Simple Linear](#)

Sponsors



Quantide: statistical consulting and training



TRAININGS: R & PYTHON
 Einführung & Machine Learning
BERLIN HAMBURG
 April 2020 Oktober 2020
[Jetzt anmelden](#)



The only thing you can be sure of is way more Data!


TRAINING: R, SCALA, STAN

STATWORX
[Data Science Service](#)
 Data Science | Consulting | Development | Training

Try the FASTEST ML for R

[Click for a Free Trial](#)

SIGMA


Our ads respect your privacy. Read our [Privacy Policy page](#) to learn more.

[Contact us](#) if you wish to help support R-bloggers, and place **your banner here**.

[Jobs for R users](#)

- [\(Junior\) Data Analyst – Berlin](#)
- [Associate Researcher \(Data Analytics\)](#)
- [Lead Data Scientist @ Stabio, Ticino, Switzerland](#)
- [Senior Scientist, Translational Informatics](#)
- [Senior Scientist, Translational Informatics @ Vancouver, BC, Canada](#)

States

- [Technical Research Analyst – New York, U.S.](#)

Full list of contributing R-bloggers

[R-bloggers](#) was founded by [Tal Galili](#), with gratitude to the [R](#) community.

Is powered by [WordPress](#) using a [bavotasan.com](#) design.

Copyright © 2019 **R-bloggers**. All Rights Reserved. [Terms and Conditions](#) for this website